

1. In this problem we have to predict the Stage of Breast Cancer M (Malignant) and B (Benign)

2. Attribute Information:

1) ID number

2) Diagnosis (M = malignant, B = benign)

Ten real-valued features are computed for each cell nucleus:

a) radius (mean of distances from center to points on the perimeter)

b) texture (standard deviation of gray-scale values)

c) perimeter

d) area

e) smoothness (local variation in radius lengths)

f) compactness ($\text{perimeter}^2 / \text{area} - 1.0$)

g) concavity (severity of concave portions of the contour)

h) concave points (number of concave portions of the contour)

i) symmetry

j) fractal dimension ("coastline approximation" - 1)

3. radius, texture, area, perimeter, smoothness, compactness, concavity, concave points, symmetry and fractal dimension these are the parameters which are very useful

In [126]:

```
import numpy as np
import pandas as pd
```

In [127]:

```
data = pd.read_csv(r"C:\Users\kotha\Downloads\cancer.csv", header=0)
```

In [128]:

```
print(data.head(2))
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	\
0	842302	M	17.99	10.38	122.8	1001.0	
1	842517	M	20.57	17.77	132.9	1326.0	

	smoothness_mean	compactness_mean	concavity_mean	concave	points_mean	\
0	0.11840	0.27760	0.3001		0.14710	
1	0.08474	0.07864	0.0869		0.07017	

	...	radius_worst	texture_worst	perimeter_worst	area_worst	\
0	...	25.38	17.33	184.6	2019.0	
1	...	24.99	23.41	158.8	1956.0	

	smoothness_worst	compactness_worst	concavity_worst	concave	points_worst
t \					
0	0.1622	0.6656	0.7119		0.265
4					
1	0.1238	0.1866	0.2416		0.186
0					

	symmetry_worst	fractal_dimension_worst
0	0.4601	0.11890
1	0.2750	0.08902

```
[2 rows x 32 columns]
```

In [129]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     569 non-null    int64
1   diagnosis                             569 non-null    object
2   radius_mean                           569 non-null    float64
3   texture_mean                           569 non-null    float64
4   perimeter_mean                         569 non-null    float64
5   area_mean                             569 non-null    float64
6   smoothness_mean                       569 non-null    float64
7   compactness_mean                      569 non-null    float64
8   concavity_mean                        569 non-null    float64
9   concave points_mean                   569 non-null    float64
10  symmetry_mean                         569 non-null    float64
11  fractal_dimension_mean                569 non-null    float64
12  radius_se                             569 non-null    float64
13  texture_se                             569 non-null    float64
14  perimeter_se                           569 non-null    float64
15  area_se                               569 non-null    float64
16  smoothness_se                         569 non-null    float64
17  compactness_se                        569 non-null    float64
18  concavity_se                          569 non-null    float64
19  concave points_se                     569 non-null    float64
20  symmetry_se                           569 non-null    float64
21  fractal_dimension_se                  569 non-null    float64
22  radius_worst                          569 non-null    float64
23  texture_worst                         569 non-null    float64
24  perimeter_worst                       569 non-null    float64
25  area_worst                            569 non-null    float64
26  smoothness_worst                     569 non-null    float64
27  compactness_worst                     569 non-null    float64
28  concavity_worst                       569 non-null    float64
29  concave points_worst                  569 non-null    float64
30  symmetry_worst                        569 non-null    float64
31  fractal_dimension_worst                569 non-null    float64
dtypes: float64(30), int64(1), object(1)
memory usage: 142.4+ KB
```

In [130]:

```
data.drop("id",axis=1,inplace=True)
```

In [131]:

```
data.isnull().any()
```

Out[131]:

```
diagnosis           False
radius_mean         False
texture_mean        False
perimeter_mean      False
area_mean           False
smoothness_mean     False
compactness_mean    False
concavity_mean      False
concave points_mean False
symmetry_mean       False
fractal_dimension_mean False
radius_se           False
texture_se          False
perimeter_se        False
area_se             False
smoothness_se       False
compactness_se      False
concavity_se        False
concave points_se   False
symmetry_se         False
fractal_dimension_se False
radius_worst        False
texture_worst       False
perimeter_worst     False
area_worst          False
smoothness_worst    False
compactness_worst   False
concavity_worst     False
concave points_worst False
symmetry_worst      False
fractal_dimension_worst False
dtype: bool
```

In [132]:

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
data['diagnosis']=le.fit_transform(data['diagnosis'])
```

In [133]:

data.head(5)

Out[133]:

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	com
0	1	17.99	10.38	122.80	1001.0	0.11840	
1	1	20.57	17.77	132.90	1326.0	0.08474	
2	1	19.69	21.25	130.00	1203.0	0.10960	
3	1	11.42	20.38	77.58	386.1	0.14250	
4	1	20.29	14.34	135.10	1297.0	0.10030	

5 rows × 31 columns

In [134]:

data.shape

Out[134]:

(569, 31)

In [135]:

data.describe()

Out[135]:

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	com
count	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	
mean	0.372583	14.127292	19.289649	91.969033	654.889104	0.096360	
std	0.483918	3.524049	4.301036	24.298981	351.914129	0.014064	
min	0.000000	6.981000	9.710000	43.790000	143.500000	0.052630	
25%	0.000000	11.700000	16.170000	75.170000	420.300000	0.086370	
50%	0.000000	13.370000	18.840000	86.240000	551.100000	0.095870	
75%	1.000000	15.780000	21.800000	104.100000	782.700000	0.105300	
max	1.000000	28.110000	39.280000	188.500000	2501.000000	0.163400	

8 rows × 31 columns

In [136]:

```
x=data.iloc[:,1:31].values  
y=data.iloc[:,0:1].values
```

In [137]:

```
x.shape
```

Out[137]:

```
(569, 30)
```

In [138]:

```
from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=0)
```

In [139]:

```
x_train.shape
```

Out[139]:

```
(455, 30)
```

In [140]:

```
x_train
```

Out[140]:

```
array([[1.005e+01, 1.753e+01, 6.441e+01, ..., 6.499e-02, 2.894e-01,  
        7.664e-02],  
       [1.080e+01, 2.198e+01, 6.879e+01, ..., 7.485e-02, 2.965e-01,  
        7.662e-02],  
       [1.614e+01, 1.486e+01, 1.043e+02, ..., 1.129e-01, 2.778e-01,  
        7.012e-02],  
       ...,  
       [9.436e+00, 1.832e+01, 5.982e+01, ..., 5.052e-02, 2.454e-01,  
        8.136e-02],  
       [9.720e+00, 1.822e+01, 6.073e+01, ..., 0.000e+00, 1.909e-01,  
        6.559e-02],  
       [1.151e+01, 2.393e+01, 7.452e+01, ..., 9.653e-02, 2.112e-01,  
        8.732e-02]])
```

In [141]:

```
from sklearn.preprocessing import StandardScaler  
sc=StandardScaler()  
x_train=sc.fit_transform(x_train)  
x_test=sc.fit_transform(x_test)
```

In [142]:

x_train

Out[142]:

```
array([[ -1.15036482, -0.39064196, -1.12855021, ..., -0.75798367,
        -0.01614761, -0.38503402],
       [ -0.93798972,  0.68051405, -0.94820146, ..., -0.60687023,
         0.09669004, -0.38615797],
       [  0.574121   , -1.03333557,  0.51394098, ..., -0.02371948,
        -0.20050207, -0.75144254],
       ...,
       [ -1.32422924, -0.20048168, -1.31754581, ..., -0.97974953,
        -0.71542314, -0.11978123],
       [ -1.24380987, -0.2245526  , -1.28007609, ..., -1.75401433,
        -1.58157125, -1.00601779],
       [ -0.73694129,  1.14989702, -0.71226578, ..., -0.27460457,
        -1.25895095,  0.21515662]])
```

In [143]:

```
from sklearn.linear_model import LogisticRegression
logistic=LogisticRegression()
logistic.fit(x_train,y_train)
```

C:\Users\kotha\anaconda3\lib\site-packages\sklearn\utils\validation.py:760:
DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

Out[143]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                    warm_start=False)
```

In [144]:

```
lpred=logistic.predict(x_test)
```

In [145]:

lpred

Out[145]:

```
array([1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 1, 1,
        0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0,
        0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0,
        1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0,
        1, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,
        0, 1, 1, 0])
```

In [146]:

```
y_test
```

Out[146]:

```
array([[1],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [1],
       [0],
       [1]])
```

In [147]:

```
from sklearn.metrics import accuracy_score
laccuracy=accuracy_score(y_test,lpred)
```

In [148]:

```
laccuracy
```

Out[148]:

```
0.956140350877193
```

In [149]:

```
from sklearn.metrics import confusion_matrix
lcm=confusion_matrix(y_test,lpred)
```

In [150]:

```
lcm
```

Out[150]:

```
array([[65,  2],
       [ 3, 44]], dtype=int64)
```

In [151]:

```
x_test.shape
```

Out[151]:

```
(114, 30)
```


In [152]:

```
import sklearn.metrics as metrics
lfr,ltpr,lthreshold=metrics.roc_curve(y_test,lpred)
lroc_auc=metrics.auc(lfr,ltpr)
```

In [153]:

```
lroc_auc
```

Out[153]:


```
0.9531597332486503
```

In [154]:

```
import matplotlib.pyplot as plt
plt.title("Receiver Opearting Characteristics")
plt.plot(lfr,ltpr,color="pink",label='Auc_roc curve value=%0.2f'%lroc_auc)
plt.legend()
plt.xlim([0,1])
plt.ylim([0,1])
plt.ylabel("true positive rate")
plt.xlabel("false positive rate")
```

Out[154]:

```
Text(0.5, 0, 'false positive rate')
```



In [159]:

```
daccuracy=accuracy_score(y_test,dpred)
daccuracy
```

Out[159]:

0.9385964912280702

In [160]:

```
dcm=confusion_matrix(y_test,dpred)
```

In [161]:

```
dcm
```

Out[161]:

```
array([[64,  3],
       [ 4, 43]], dtype=int64)
```

In [210]:

```
import sklearn.metrics as metrics
dfpr,dtp,dthreshold=metrics.roc_curve(y_test,dpred)
droc_auc=metrics.auc(dfpr,dtp)
```

In [211]:

```
droc_auc
```

Out[211]:


0.9350587488091456

In [163]:

```
import matplotlib.pyplot as plt
plt.title("Receiver Operating Characteristics")
plt.plot(dfpr,dtptr,color="orange",label='Auc_roc curve value=%0.2f'%droc_auc)
plt.legend()
plt.xlim([0,1])
plt.ylim([0,1])
plt.ylabel("true positive rate")
plt.xlabel("false positive rate")
```

Out[163]:

Text(0.5, 0, 'false positive rate')



In [164]:

```
from sklearn.ensemble import RandomForestClassifier
random=RandomForestClassifier(n_estimators=10,criterion="entropy",random_state=0)
random.fit(x_train,y_train)
```

C:\Users\kotha\anaconda3\lib\site-packages\ipykernel_launcher.py:3: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

This is separate from the ipykernel package so we can avoid doing imports until

Out[164]:

```
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                        criterion='entropy', max_depth=None, max_features='au
to',
                        max_leaf_nodes=None, max_samples=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators=10,
                        n_jobs=None, oob_score=False, random_state=0, verbose
=0,
                        warm_start=False)
```

In [165]:

```
rpred=random.predict(x_test)
```

In [166]:

```
rpred
```

Out[166]:

```
array([1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1,
       0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0,
       0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0,
       1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0,
       1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1,
       0, 1, 1, 0])
```

In [167]:

```
y_test
```

Out[167]:

```
array([[1],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [0],
       [1],
       [0],
       [1].
```

In [168]:

```
raccuracy=accuracy_score(y_test, rpred)
raccuracy
```

Out[168]:

```
0.9736842105263158
```

In [169]:

```
rcm=confusion_matrix(y_test, rpred)
```

In [170]:

```
rcm
```

Out[170]:

```
array([[66,  1],  
       [ 2, 45]], dtype=int64)
```

In [171]:

```
import sklearn.metrics as metrics  
rfpr,rtp,rthreshold=metrics.roc_curve(y_test,rpred)  
rroc_auc=metrics.auc(rfpr,rtp)
```

In [172]:

```
rroc_auc
```

Out[172]:


```
0.971260717688155
```

In [212]:

```
import matplotlib.pyplot as plt  
plt.title("Receiver Opearting Characteristics")  
plt.plot(dfpr,dtp,color="purple",label='Auc_roc curve value=%0.2f'%rroc_auc)  
plt.legend()  
plt.xlim([0,1])  
plt.ylim([0,1])  
plt.ylabel("true positive rate")  
plt.xlabel("false positive rate")
```

Out[212]:

```
Text(0.5, 0, 'false positive rate')
```



```
from sklearn.neighbors import KNeighborsClassifier  
knn= KNeighborsClassifier(n_neighbors=5,metric="minkowski",p=2)  
knn.fit(x_train,y_train)
```


In [181]:

```
kcm
```

Out[181]:

```
array([[67,  0],  
       [ 4, 43]], dtype=int64)
```

In [182]:

```
import sklearn.metrics as metrics  
kfpr,ktpr,kthreshold=metrics.roc_curve(y_test,kpred)  
kroc_auc=metrics.auc(kfpr,ktpr)
```

In [183]:

```
kroc_auc
```

Out[183]:


```
0.9574468085106382
```

In [184]:

```
import matplotlib.pyplot as plt  
plt.title("Receiver Opearting Characteristics")  
plt.plot(kfpr,ktpr,color="blue",label='Auc_roc curve value=%0.2f'%kroc_auc)  
plt.legend()  
plt.xlim([0,1])  
plt.ylim([0,1])  
plt.ylabel("true positive rate")  
plt.xlabel("false positive rate")
```

Out[184]:

```
Text(0.5, 0, 'false positive rate')
```



In [185]:

ktp

Out[185]:

array([0. , 0.91489362, 1.])

In [186]:

kfpr

Out[186]:

array([0., 0., 1.])

In [187]:

kthreshold

Out[187]:

array([2, 1, 0])

In [188]:

```
from sklearn.naive_bayes import GaussianNB
naive=GaussianNB()
naive.fit(x_train,y_train)
```

C:\Users\kotha\anaconda3\lib\site-packages\sklearn\naive_bayes.py:206: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

Out[188]:

GaussianNB(priors=None, var_smoothing=1e-09)

In [189]:

npred=naive.predict(x_test)

In [190]:

npred

Out[190]:

```
array([1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 1,
       0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0,
       0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0,
       1, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0,
       1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1,
       0, 1, 1, 0])
```

In [191]:

```
y_test
```

Out[191]:

```
array([[1],  
       [0],  
       [0],  
       [0],  
       [0],  
       [0],  
       [0],  
       [0],  
       [0],  
       [0],  
       [0],  
       [0],  
       [0],  
       [0],  
       [0],  
       [1],  
       [0],  
       [1]])
```

In [192]:

```
naccuracy=accuracy_score(y_test,npred)
```

In [193]:

```
naccuracy
```

Out[193]:

```
0.9385964912280702
```

In [194]:

```
ncm=confusion_matrix(y_test,npred)
```

In [195]:

```
ncm
```

Out[195]:

```
array([[64,  3],  
       [ 4, 43]], dtype=int64)
```

In [196]:

```
import sklearn.metrics as metrics  
nfpr,ntpr,nthreshold=metrics.roc_curve(y_test,npred)  
nroc_auc=metrics.auc(nfpr,ntpr)
```

In [197]:

nroc_auc

Out[197]:

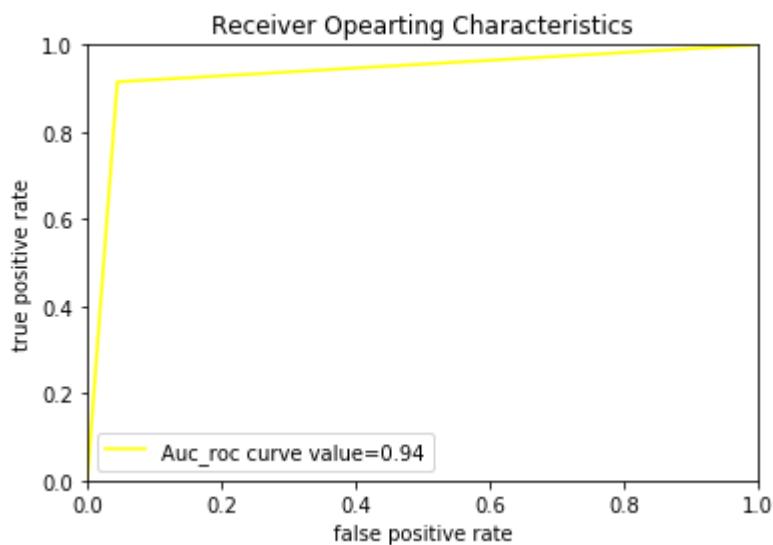
0.9350587488091456

In [198]:

```
import matplotlib.pyplot as plt
plt.title("Receiver Operating Characteristics")
plt.plot(nfpr,ntpr,color="yellow",label='Auc_roc curve value=%0.2f'%nroc_auc)
plt.legend()
plt.xlim([0,1])
plt.ylim([0,1])
plt.ylabel("true positive rate")
plt.xlabel("false positive rate")
```

Out[198]:

Text(0.5, 0, 'false positive rate')



In [199]:

```
from sklearn.svm import SVC
svm=SVC(kernel="linear")
svm.fit(x_train,y_train)
```

C:\Users\kotha\anaconda3\lib\site-packages\sklearn\utils\validation.py:760:
DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, warn=True)

Out[199]:

```
SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,  
    decision_function_shape='ovr', degree=3, gamma='scale', kernel='linear',  
    max_iter=-1, probability=False, random_state=None, shrinking=True,  
    tol=0.001, verbose=False)
```

```
spred=svm.predict(x_test)
```

spread

```
array([1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1,
       0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0,
       0, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0,
       1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0,
       1, 1, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,
       0, 1, 1, 0])
```

y_test

[illegible]

```
saccuracy=accuracy_score(y_test,spred)
```

saccuracy

0.9736842105263158

```
scm=confusion_matrix(y_test,spred)
```

In [206]:

scm

Out[206]:

```
array([[66,  1],
       [ 2, 45]], dtype=int64)
```

In [207]:

```
import sklearn.metrics as metrics
sfpr,stpr,sthreshold=metrics.roc_curve(y_test,spred)
sroc_auc=metrics.auc(sfpr,stpr)
```

In [208]:

sroc_auc

Out[208]:

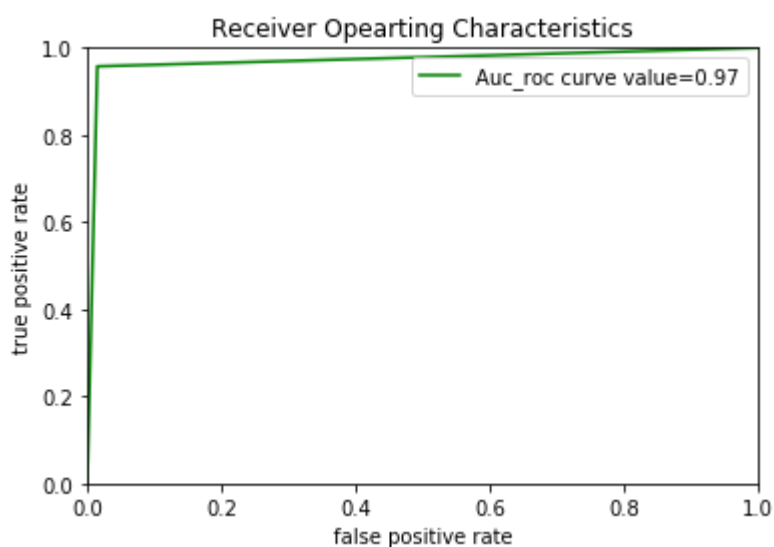
0.971260717688155

In [209]:

```
import matplotlib.pyplot as plt
plt.title("Receiver Opearting Characteristics")
plt.plot(sfpr,stpr,color="green",label='Auc_roc curve value=%0.2f'%sroc_auc)
plt.legend()
plt.xlim([0,1])
plt.ylim([0,1])
plt.ylabel("true positive rate")
plt.xlabel("false positive rate")
```

Out[209]:

Text(0.5, 0, 'false positive rate')



Observation:

LOGISTIC:

NAIVE_BAYES:

laccuracy=0.95

naccuracy=0.93

DECISIONTREE:

daccuracy=0.938

RANDOMFOREST:

raccuracy=0.97

KNN:

kaccuracy=0.96

```
Auc_roc curve value=0.95   Auc_roc curve value=0.94   Auc_roc value=0.97   Auc_roc curve  
value=0.96 Auc_roc curve value=0.94
```

SVM:

saccuracy=0.97

Auc_roc curve value=0.97

In []: